

# Stats Homework 2

## Environment setup

Same initial steps that we ran in Stats Demo 1, just consolidated.

```
knitr::opts_chunk$set(echo = TRUE, tidy.opts = list(width.cutoff = 60), tidy = TRUE)

library(Sleuth3) # example datasets from textbook, "The Statistical Sleuth - A Course in
# Methods of Data Analysis (3rd Edition)"
library(MASS) # example dataset representing the sales of different models of cars in 1993
library(reshape2) # for formatting and aggregation of data frames
library(ggplot2) # for creating graphs
library(dplyr) # for data manipulation and clean-up
library(plotly) # for creating interactive web graphics from ggplot2 graphs
library(knitr) # required for generating PDF output
library(modeest) # required for `mfu()` function
library(nortest) # required for running Anderson-Darling test for normality
library(effsize)
```

## ANOVA

As an SE researcher you are evaluating different programming languages. For the next set of questions input the R code and interpret your findings.

- a) The results of your first study compares Java, Python, and Ruby code based on the size of the programs in source (i.e. non-blank, non-commented) lines of code. Perform an ANOVA to determine whether there is an effect on size due to programming language. Use `lang-size.csv`.

```
# code goes here.
data_a <- read.csv("lang-size.csv")
data_a
```

```
##      lang sloc
## 1    java 207
## 2    java 296
## 3    java 348
## 4    java 309
## 5    java 231
## 6    java 228
## 7    java 318
## 8    java 212
## 9    java 284
## 10   java 267
## 11   java 354
```

##	12	java	262
##	13	java	259
##	14	java	342
##	15	java	252
##	16	java	299
##	17	java	312
##	18	java	285
##	19	java	266
##	20	java	333
##	21	java	280
##	22	java	283
##	23	java	368
##	24	java	407
##	25	java	325
##	26	java	339
##	27	java	255
##	28	java	327
##	29	java	268
##	30	java	315
##	31	python	401
##	32	python	391
##	33	python	331
##	34	python	370
##	35	python	415
##	36	python	419
##	37	python	340
##	38	python	382
##	39	python	319
##	40	python	387
##	41	python	455
##	42	python	438
##	43	python	388
##	44	python	449
##	45	python	437
##	46	python	404
##	47	python	352
##	48	python	390
##	49	python	446
##	50	python	424
##	51	python	370
##	52	python	291
##	53	python	366
##	54	python	294
##	55	python	337
##	56	python	381
##	57	python	366
##	58	python	356
##	59	python	395
##	60	python	387
##	61	ruby	188
##	62	ruby	227
##	63	ruby	208
##	64	ruby	267
##	65	ruby	303

```
## 66  ruby 311
## 67  ruby 287
## 68  ruby 226
## 69  ruby 278
## 70  ruby 188
## 71  ruby 269
## 72  ruby 178
## 73  ruby 198
## 74  ruby 239
## 75  ruby 176
## 76  ruby 309
## 77  ruby 176
## 78  ruby 228
## 79  ruby 197
## 80  ruby 326
## 81  ruby 280
## 82  ruby 239
## 83  ruby 286
## 84  ruby 286
## 85  ruby 272
## 86  ruby 258
## 87  ruby 283
## 88  ruby 361
## 89  ruby 191
## 90  ruby 246
```

```
summary(aov(sloc ~ lang, data = data_a))
```

```
##           Df Sum Sq Mean Sq F value Pr(>F)
## lang       2 276056   138028    62.6 <2e-16 ***
## Residuals  87 191836     2205
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Report: We ran Anova( $F(2, 87) = 62.6$ ,  $p < 0.001$ )

The p-value is lower than 0.05, and the differences are statistically significant, which is the language has a real impact.

- b) In a subsequent study you measured the programming time (in hours) required to solve a program in Java, Python, and Ruby. This was a within subject study design: each participant solved the problem three times, and all participants solved the problem in the same order (Java, then Python, then Ruby). Perform an ANOVA to determine whether there is an effect due to programming language. Use `lang-time.csv`.

```
# code goes here
data_b <- read.csv("lang-time.csv")
data_b
```

```
##      lang participant times
## 1    java          P1   11.3
## 2    java          P2    7.8
## 3    java          P3   12.6
```

## 4	java	P4	6.3
## 5	java	P5	8.1
## 6	java	P6	10.1
## 7	java	P7	3.2
## 8	java	P8	7.3
## 9	java	P9	8.1
## 10	java	P10	7.9
## 11	java	P11	9.0
## 12	java	P12	8.7
## 13	java	P13	7.7
## 14	java	P14	11.6
## 15	java	P15	6.0
## 16	java	P16	8.1
## 17	java	P17	10.9
## 18	java	P18	9.1
## 19	java	P19	8.8
## 20	java	P20	9.8
## 21	java	P21	9.5
## 22	java	P22	8.7
## 23	java	P23	10.4
## 24	java	P24	9.7
## 25	python	P1	7.5
## 26	python	P2	5.1
## 27	python	P3	6.8
## 28	python	P4	5.7
## 29	python	P5	3.1
## 30	python	P6	6.8
## 31	python	P7	7.6
## 32	python	P8	11.2
## 33	python	P9	11.1
## 34	python	P10	10.2
## 35	python	P11	5.8
## 36	python	P12	4.3
## 37	python	P13	6.2
## 38	python	P14	10.5
## 39	python	P15	8.2
## 40	python	P16	8.8
## 41	python	P17	6.4
## 42	python	P18	4.9
## 43	python	P19	9.1
## 44	python	P20	7.3
## 45	python	P21	5.9
## 46	python	P22	8.0
## 47	python	P23	9.8
## 48	python	P24	9.8
## 49	ruby	P1	9.1
## 50	ruby	P2	6.2
## 51	ruby	P3	9.2
## 52	ruby	P4	6.9
## 53	ruby	P5	7.4
## 54	ruby	P6	9.7
## 55	ruby	P7	8.5
## 56	ruby	P8	7.1
## 57	ruby	P9	9.7

```
## 58  ruby      P10  6.2
## 59  ruby      P11  5.4
## 60  ruby      P12  7.7
## 61  ruby      P13  7.1
## 62  ruby      P14  6.9
## 63  ruby      P15  6.2
## 64  ruby      P16  5.3
## 65  ruby      P17  9.4
## 66  ruby      P18  3.7
## 67  ruby      P19  9.3
## 68  ruby      P20  4.9
## 69  ruby      P21  8.8
## 70  ruby      P22  4.4
## 71  ruby      P23  8.3
## 72  ruby      P24  5.7
```

```
summary(aov(times ~ lang + participant, data = data_b))
```

```
##              Df Sum Sq Mean Sq F value Pr(>F)
## lang           2  33.32   16.661    4.583 0.0153 *
## participant    23 110.62    4.809    1.323 0.2061
## Residuals     46  167.24    3.636
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Report: We ran Anova ( $F(2, 46) = 4.583$ ,  $p = 0.0153$ ) , Anova ( $F(23, 46) = 1.323$ ,  $p = 0.2061$ )

The p-value of lang is lower than 0.05 showing substantial evidence against the null hypothesis.

- c) You realized you should have counterbalanced, so you replicated the study from (b) which uses a crossover design to control for ordering. Each participant solved the problem in all three languages, but in each participant solved them in a different order. Perform an ANOVA to determine whether there is an effect due to programming language. Use `lang-time-crossover.csv`.

```
# code goes here
data_c <- read.csv("lang-time-crossover.csv")
data_c
```

```
## participant treatment lang times
## 1          P1         T1  java  6.4
## 2          P2         T1  java  8.3
## 3          P3         T1 python 7.0
## 4          P4         T1 python 10.5
## 5          P5         T1  ruby 10.6
## 6          P6         T1  ruby  4.0
## 7          P1         T2 python 8.2
## 8          P2         T2  ruby  5.5
## 9          P3         T2  java  7.7
## 10         P4         T2  ruby  7.5
## 11         P5         T2  java  7.0
## 12         P6         T2 python 4.4
## 13         P1         T3  ruby  5.7
## 14         P2         T3 python 7.9
```

```
## 15      P3      T3    ruby    8.8
## 16      P4      T3     java    9.5
## 17      P5      T3  python    8.0
## 18      P6      T3     java    8.0
```

```
summary(aov(times ~ lang + participant + treatment, data = data_c))
```

```
##           Df Sum Sq Mean Sq F value Pr(>F)
## lang       2  2.170   1.085   0.377  0.698
## participant 5 26.100   5.220   1.812  0.217
## treatment   2  5.623   2.812   0.976  0.418
## Residuals   8 23.047   2.881
```

Report: We ran Anova ( $F(2, 8) = 0.377$ ,  $p = 0.698$ ) , Anova ( $F(5, 8) = 1.812$ ,  $p = 0.217$ ), Anova ( $F(2, 8) = 0.976$ ,  $p = 0.418$ )

The p-value of lang is higher than 0.05, there's no strong evidence against null hypothesis. We cannot conclude that a significant difference exists.

- d) You have some simulated results from an experiment that compared development time for Java, Python and Ruby, for subjects with low experience and high experience. Perform an ANOVA and identify which factors (language, experience) had a statistically significant effect. Also specify whether the interaction between programming language and experience was statistically significant or not. Use `lang-time-exp.csv`. yes

```
# code goes here
data_d <- read.csv("lang-time-exp.csv")
data_d
```

```
##      lang exp times
## 1    java low  11.0
## 2    java low  10.6
## 3    java low   8.3
## 4    java low   9.8
## 5    java low  11.6
## 6    java low  11.8
## 7    java low   8.6
## 8    java low  10.3
## 9    java low   7.7
## 10   java low  10.5
## 11   java low  13.2
## 12   java low  12.5
## 13   java low  10.5
## 14   java low  13.0
## 15   java low  12.5
## 16   java low  11.2
## 17   java low   9.1
## 18   java low  10.6
## 19   java low  12.9
## 20   java low  12.0
## 21   java high  5.8
## 22   java high  2.6
## 23   java high  5.7
```

## 24	java	high	2.8
## 25	java	high	4.5
## 26	java	high	6.3
## 27	java	high	5.6
## 28	java	high	5.3
## 29	java	high	6.8
## 30	java	high	6.5
## 31	java	high	3.3
## 32	java	high	6.8
## 33	java	high	8.9
## 34	java	high	7.4
## 35	java	high	4.2
## 36	java	high	4.1
## 37	java	high	7.7
## 38	java	high	3.5
## 39	java	high	6.4
## 40	java	high	5.7
## 41	python	low	12.2
## 42	python	low	8.5
## 43	python	low	8.3
## 44	python	low	11.7
## 45	python	low	8.1
## 46	python	low	9.9
## 47	python	low	10.5
## 48	python	low	9.4
## 49	python	low	8.6
## 50	python	low	11.3
## 51	python	low	9.2
## 52	python	low	9.3
## 53	python	low	12.7
## 54	python	low	14.3
## 55	python	low	11.0
## 56	python	low	11.6
## 57	python	low	8.2
## 58	python	low	11.1
## 59	python	low	8.7
## 60	python	low	10.6
## 61	python	high	3.5
## 62	python	high	5.1
## 63	python	high	4.3
## 64	python	high	6.7
## 65	python	high	8.1
## 66	python	high	8.4
## 67	python	high	7.5
## 68	python	high	5.0
## 69	python	high	7.1
## 70	python	high	3.5
## 71	python	high	6.8
## 72	python	high	3.1
## 73	python	high	3.9
## 74	python	high	5.6
## 75	python	high	3.0
## 76	python	high	8.3
## 77	python	high	3.0

```
## 78 python high 5.1
## 79 python high 3.9
## 80 python high 9.0
## 81 ruby low 10.2
## 82 ruby low 8.6
## 83 ruby low 10.4
## 84 ruby low 10.4
## 85 ruby low 9.9
## 86 ruby low 9.3
## 87 ruby low 10.3
## 88 ruby low 13.4
## 89 ruby low 6.6
## 90 ruby low 8.9
## 91 ruby low 8.2
## 92 ruby low 8.6
## 93 ruby low 9.1
## 94 ruby low 11.3
## 95 ruby low 10.2
## 96 ruby low 6.2
## 97 ruby low 5.8
## 98 ruby low 10.8
## 99 ruby low 9.3
## 100 ruby low 11.5
## 101 ruby high 3.5
## 102 ruby high 5.8
## 103 ruby high 2.9
## 104 ruby high 6.4
## 105 ruby high 3.7
## 106 ruby high 6.1
## 107 ruby high 6.3
## 108 ruby high 1.7
## 109 ruby high 7.1
## 110 ruby high 7.3
## 111 ruby high 2.4
## 112 ruby high 7.3
## 113 ruby high 4.1
## 114 ruby high 4.4
## 115 ruby high 6.8
## 116 ruby high 5.2
## 117 ruby high 7.1
## 118 ruby high 6.5
## 119 ruby high 7.5
## 120 ruby high 6.9
```

```
summary(aov(times ~ lang + exp + lang:exp, data = data_d))
```

```
##           Df Sum Sq Mean Sq F value Pr(>F)
## lang       2   11.1     5.6   1.730  0.182
## exp        1  663.2   663.2  206.137 <2e-16 ***
## lang:exp    2    9.7     4.8   1.502  0.227
## Residuals 114  366.8     3.2
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```



Report: We ran Anova ( $F(2, 114) = 1.730, p = 0.182$ ), Anova ( $F(1, 114) = 206.137, p < 0.001$ ), Anova ( $F(2, 114) = 1.502, p = 0.227$ )

The exp p-value is lower than 0.05, which means the differences are statistically significant.

The p-value between language and experience is 0.227; we cannot conclude that a significant difference exists.

## Part 3: Data analysis of an experiment

In this question, you'll analyze the raw data from an experiment and write up the results (similar to a publication).

The data is from an experiment to test whether statically typed languages (e.g. Java) or dynamically typed languages (e.g. Python) require more programming effort. The study evaluates the languages on two problems, a "small" problem and a "large" problem, to see if the results change based on the size of the problem. The study is a factorial design. The raw data from the experiment is available in this file: `lang-time-size.csv`.

Analyze the data and write up a short "results" section (as if it were a part of a paper) with your analysis of the data. This section should contain: \* Box plots to show the raw data \* Analysis of variance tables to determine if there are any interactions \* Results of pairwise t-tests to test if the factors that you may see significance truly have an effect \* Interaction plot between the 2 factors \* Effect sizes for programming language for the "small" problem and for the "large" problem. \* I am not looking for a specific format, use your judgement about the best way to present this data to convey the results to a reader.

## NEW

## Part 3: Data analysis of an experiment

In this question, you'll analyze the raw data from an experiment and write up the results (similar to a publication).

The data is from an experiment to test whether statically typed languages (e.g. Java) or dynamically typed languages (e.g. Python) require more programming effort. The study evaluates the languages on two problems, a "small" problem and a "large" problem, to see if the results change based on the size of the problem. The study is a factorial design. The raw data from the experiment is available in this file: `lang-time-size.csv`.

Analyze the data and write up a short "results" section (as if it were a part of a paper) with your analysis of the data. This section should contain: \* Analysis of variance tables to determine if there are any interactions \* Interaction plot between the 2 factors \* Effect sizes for programming language for the "small" problem and for the "large" problem. \* I am not looking for a specific format, use your judgement about the best way to present this data to convey the results to a reader.

```
# Code for analysis goes here.
```

```
data_part3 <- read.csv("lang-time-size.csv")
data_part3
```

```
##      times  lang  size
## 1    14.0   java small
## 2    20.3   java small
## 3    11.6   java small
## 4    16.6   java small
## 5    15.0   java small
```

## 6	8.0	java small
## 7	13.1	java small
## 8	17.4	java small
## 9	12.5	java small
## 10	8.7	java small
## 11	16.4	java small
## 12	11.5	java small
## 13	13.7	java small
## 14	8.0	java small
## 15	18.8	java small
## 16	13.0	java small
## 17	12.9	java small
## 18	13.1	java small
## 19	8.3	java small
## 20	10.9	java small
## 21	18.5	java small
## 22	18.6	java small
## 23	11.4	java small
## 24	11.2	java small
## 25	17.0	java small
## 26	14.7	java small
## 27	15.9	java small
## 28	10.3	java small
## 29	14.3	java small
## 30	16.6	java small
## 31	15.3	java large
## 32	19.0	java large
## 33	30.6	java large
## 34	25.0	java large
## 35	26.7	java large
## 36	22.7	java large
## 37	17.1	java large
## 38	27.3	java large
## 39	13.9	java large
## 40	8.5	java large
## 41	21.9	java large
## 42	24.8	java large
## 43	38.9	java large
## 44	11.1	java large
## 45	16.3	java large
## 46	29.0	java large
## 47	10.0	java large
## 48	31.5	java large
## 49	24.7	java large
## 50	26.5	java large
## 51	24.5	java large
## 52	27.2	java large
## 53	17.1	java large
## 54	32.4	java large
## 55	22.3	java large
## 56	11.9	java large
## 57	13.4	java large
## 58	23.2	java large
## 59	28.3	java large

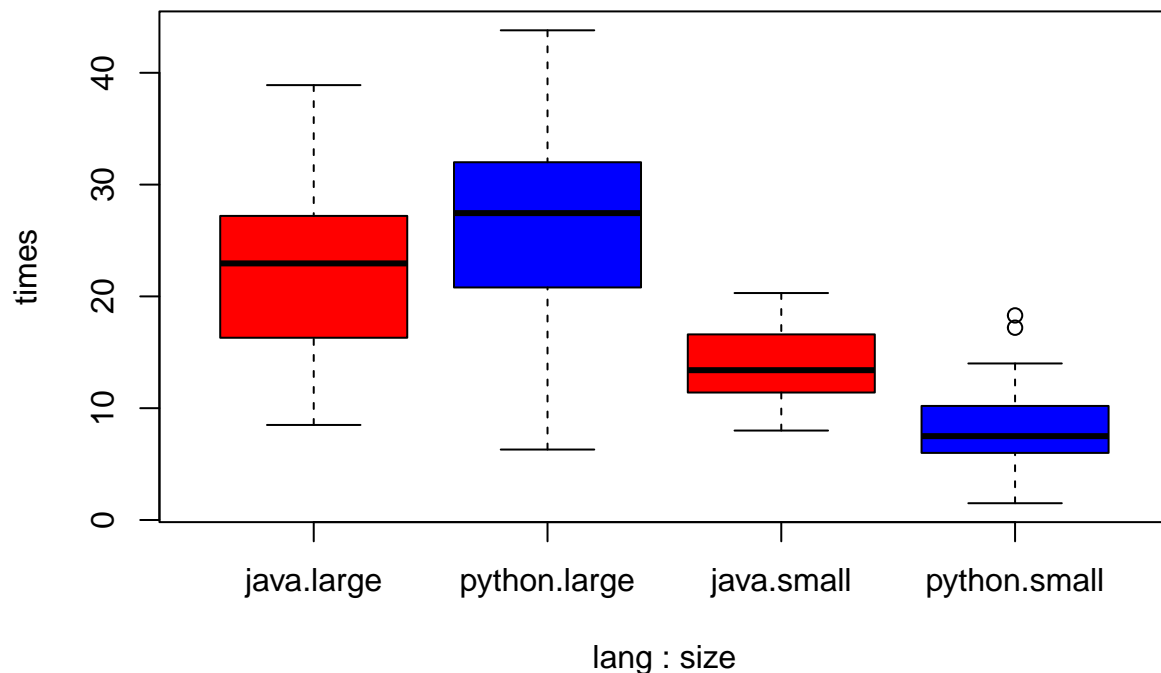
```
## 60 18.1 java large
## 61 10.2 python small
## 62 7.3 python small
## 63 6.6 python small
## 64 3.8 python small
## 65 3.0 python small
## 66 4.4 python small
## 67 7.7 python small
## 68 11.2 python small
## 69 6.0 python small
## 70 14.0 python small
## 71 2.9 python small
## 72 1.5 python small
## 73 3.7 python small
## 74 7.0 python small
## 75 9.2 python small
## 76 6.8 python small
## 77 10.4 python small
## 78 13.0 python small
## 79 5.1 python small
## 80 17.2 python small
## 81 9.0 python small
## 82 18.3 python small
## 83 9.6 python small
## 84 7.1 python small
## 85 7.1 python small
## 86 7.0 python small
## 87 7.7 python small
## 88 11.5 python small
## 89 9.5 python small
## 90 8.4 python small
## 91 20.8 python large
## 92 16.5 python large
## 93 43.8 python large
## 94 30.0 python large
## 95 25.2 python large
## 96 39.4 python large
## 97 25.6 python large
## 98 37.3 python large
## 99 19.0 python large
## 100 7.5 python large
## 101 26.7 python large
## 102 26.8 python large
## 103 17.3 python large
## 104 38.7 python large
## 105 35.6 python large
## 106 24.3 python large
## 107 28.6 python large
## 108 34.4 python large
## 109 6.3 python large
## 110 30.3 python large
## 111 20.0 python large
## 112 32.0 python large
## 113 28.1 python large
```

```
## 114 30.2 python large
## 115 29.7 python large
## 116 31.9 python large
## 117 14.9 python large
## 118 23.3 python large
## 119 35.4 python large
## 120 25.4 python large
```

```
summary(aov(times ~ lang + size + lang:size, data = data_part3))
```

```
##           Df Sum Sq Mean Sq F value    Pr(>F)
## lang       1      3      3    0.085    0.772
## size      1   5410    5410  133.246 < 2e-16 ***
## lang:size  1    811     811   19.968 1.84e-05 ***
## Residuals 116   4709      41
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
boxplot(times ~ lang + size, data = data_part3, col=c("red", "blue"))
```

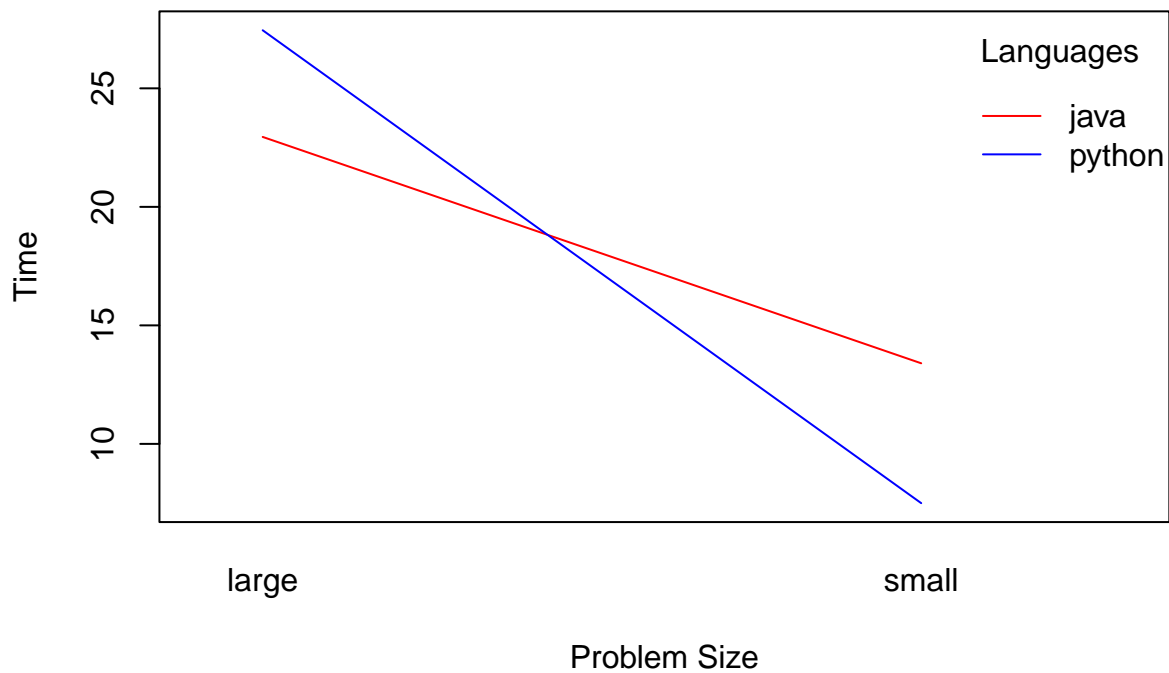


```
interaction.plot(x.factor = data_part3$size, #x-axis variable
                 trace.factor = data_part3$lang, #variable for lines
                 response = data_part3$times, #y-axis variable
                 fun = median, #metric to plot)
```

```

ylab = "Time",
xlab = "Problem Size",
col = c("red", "blue"),
lty = 1, #line type
lwd = 1, #line width
trace.label = "Languages")

```



```
# Effect sizes
```

```

small_java <- data_part3[which(data_part3$size == "small" & data_part3$lang == "java"),]
small_python <- data_part3[which(data_part3$size == "small" & data_part3$lang == "python"),]

large_java <- data_part3[which(data_part3$size == "large" & data_part3$lang == "java"),]
large_python <- data_part3[which(data_part3$size == "large" & data_part3$lang == "python"),]

cat("small problem effect sizes:\n")

```

```
## small problem effect sizes:
```

```
cohen.d(small_java$times, small_python$times)
```

```

##
## Cohen's d
##

```

```
## d estimate: 1.500909 (large)
## 95 percent confidence interval:
##      lower      upper
## 0.9158068 2.0860117
```

```
cat("larg problem effect sizes:\n")
```

```
## larg problem effect sizes:
```

```
cohen.d(large_java$times, large_python$times)
```

```
##
## Cohen's d
##
## d estimate: -0.5911475 (medium)
## 95 percent confidence interval:
##      lower      upper
## -1.11915639 -0.06313861
```

```
#cohen.d(size$java, size$python)
```

## part results: Analysis from the part3

Report: We ran Anova ( $F(1, 116) = 0.085$ ,  $p = 0.772$ ) , Anova ( $F(1, 114) = 133.246$ ,  $p < 0.001$ ), Anova ( $F(1, 116) = 19.968$ ,  $p < 0.001$ )

The p-value of size and the interactions between language and size are lower than 0.001. Therefore the differences are statistically significant. The size and the interaction between language and size have impacts on time.

By using the Cohen's function to evaluate effect size. The Cohen's d in small problem shows it has a large effect size, and the large problem doesn't.